

Talk in Parallel Computing



OpenCL

A brief introduction for
NVIDIA CUDA programmers

Jan Beneke

16.01.2012

Outline

- Motivation
- Introducing OpenCL
- Design Goals
- Hardware/ Execution Model
- Software Stack
- Known Example
- Benchmark
- Summary

Outline

- Motivation
- Introducing OpenCL
- Design Goals
- Hardware/ Execution Model
- Software Stack
- Known Example
- Benchmark
- Summary

Motivation

Advantages of GPGPU

- Runs parallelized code on GPUs
- Speeds up existing applications
- Uses ordinary and „cheap“ hardware accessible to almost every user
- Even supported by modern handheld devices
- Easy to code with dedicated language extensions and frameworks
- But: Difficult to create efficient code

Motivation

NVIDIA CUDA

Pro

- Easy to learn language extension for using GPGPU
- Intuitive framework
- Many extensions provided by NVIDIA (CUFFT, CUBLAS, ...)
- Integrated graphics language support (OpenGL, Direct 3D)
- Cross-platform (Microsoft Windows, Linux, Apple Mac OS X)
- Free of charge

Contra

- Only NVIDIA hardware supported
- Only GPUs supported
- Proprietary product/ Closed source

Outline

- Motivation
- Introducing OpenCL
- Design Goals
- Hardware/ Execution Model
- Software Stack
- Known Example
- Benchmark
- Summary

Introducing OpenCL

Open Computing Language

- Language extension for accessing heterogeneous computational devices
- Supports parallel execution on single or multiple devices of many vendors
 - GPUs (NVIDIA, AMD/ ATI, ...)
 - CPUs (Intel, AMD, ...)
 - Accelerator cards/ Server blades (IBM, ...)
- Desktop, server farm and hand-held profiles
- Integrates OpenGL and Microsoft Direct 3D graphics APIs

Introducing OpenCL

An open standard

- Vendor neutral
- Specifications under review by Khronos OpenCL working group
 - Khronos also reviews standards like OpenGL, WebGL, Open AL, ...
- Standard based on proposal by Apple and developed with industry leaders



Source: <http://www.khronos.org>

- First released and natively integrated in Mac OS X since Snow Leopard (2009)

Outline

- Motivation
- Introducing OpenCL
- Design Goals
- Hardware/ Execution Model
- Software Stack
- Known Example
- Benchmark
- Summary

Design Goals

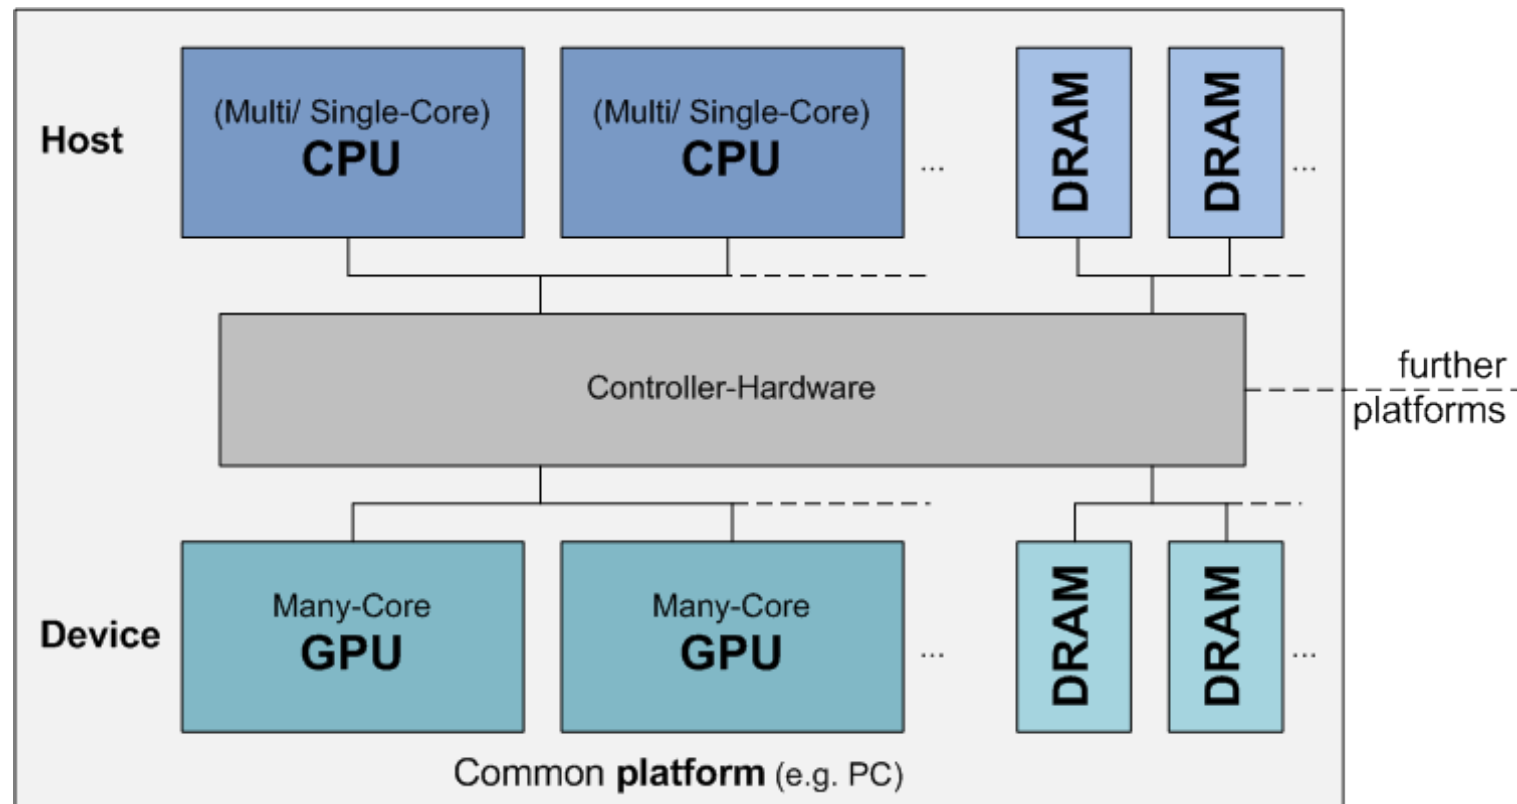
Heterogeneous World

- A modern platform includes:
 - One or more CPUs
 - One or more GPUs
 - (DSPs)
 - (Accelerator cards)
- ➔ Use them as peers.
- ➔ One portable program uses all available resources.

Design Goals

Heterogeneous World

- A common platform (PC or server blade):



Becoming a useable standard

- Efficient parallel model
 - Based on ISO C99
 - Abstract the underlying hardware
 - Additional work-items and workgroups, vector types, synchronization, address space qualifiers
 - Built-in functions for image manipulation, work-item manipulation, math routines, ...
- Specify accuracy of floating-point computations (IEEE 754)
- Support of most future hardware


Outline

- Motivation
- Introducing OpenCL
- Design Goals
- Hardware/ Execution Model
- Software Stack
- Known Example
- Benchmark
- Summary

Hardware Model

Abstract Platform Model

- One host and multiple computing devices
 - Each device contains computing units
 - Each unit is divided into processing elements



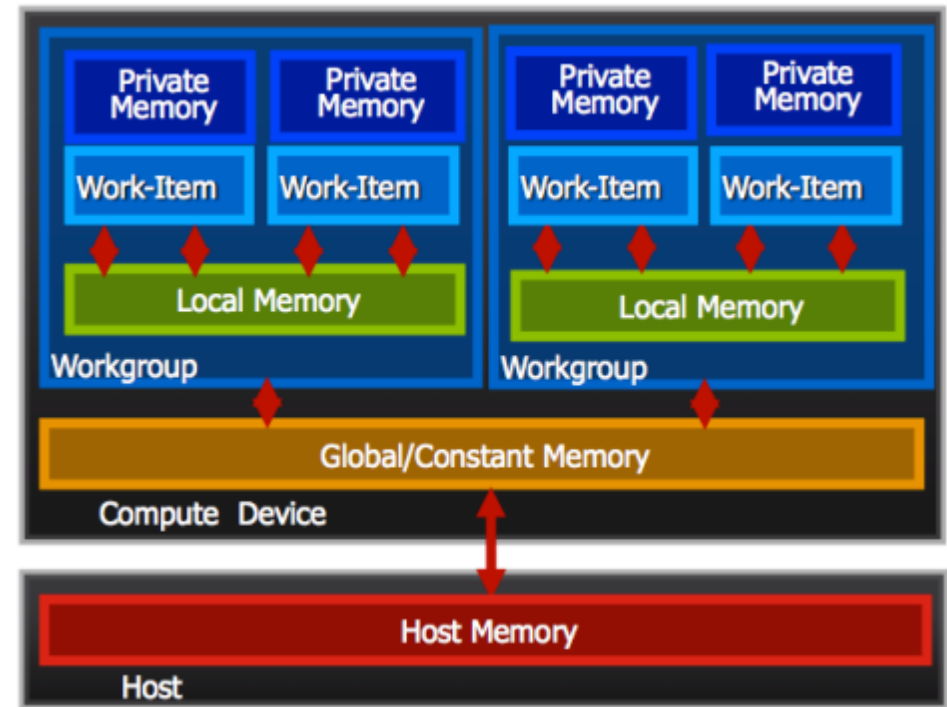
Source: <http://www.khronos.org>

Hardware Model

Abstract Memory Model

- Hardware is abstracted by a consistent shared memory model (like CUDA):

- Private Memory**
 - Per work-item
- Local Memory**
 - Shared in a workgroup
- Global/ Constant Memory**
 - All workgroups on one device
- Host Memory**
 - On host level (CPU)




Source: <http://www.khronos.org>

- Independent of based hardware.
- Explicit memory model:
 - Mostly move data from **host** → **global** → **local** → ... and back

Execution Model

The OpenCL application

- The application runs on a host, that submits work to the compute devices
 - Work-item
 - Basic unit on OpenCL device
 - Kernel
 - The code for a work item - an extended C function
 - Command Queue
 - Queues Kernels/ synchronizing - event handling
 - Program
 - Collection of kernels and other functions - analogous to a dynamic library
 - Context
 - The environment to execute the work-items in, includes devices and their memories and command queues

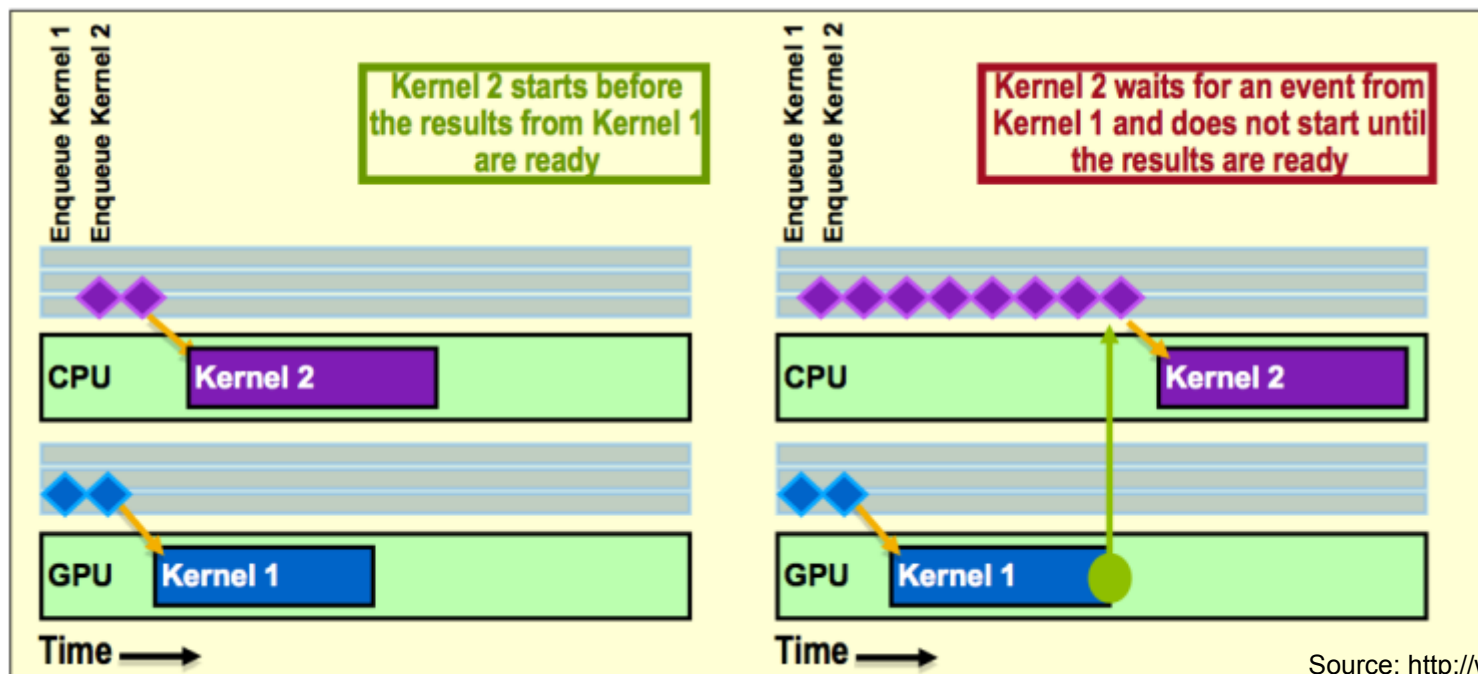


Source: <http://www.khronos.org>

Execution Model

Synchronization

- Applications queue compute kernel execution instances
- Events can be used to synchronize kernel executions between queues
- Example: 2 queues on separate devices



Source: <http://www.khronos.org>

Outline

- Motivation
- Introducing OpenCL
- Design Goals
- Hardware/ Execution Model
- Software Stack
- Known Example
- Benchmark
- Summary

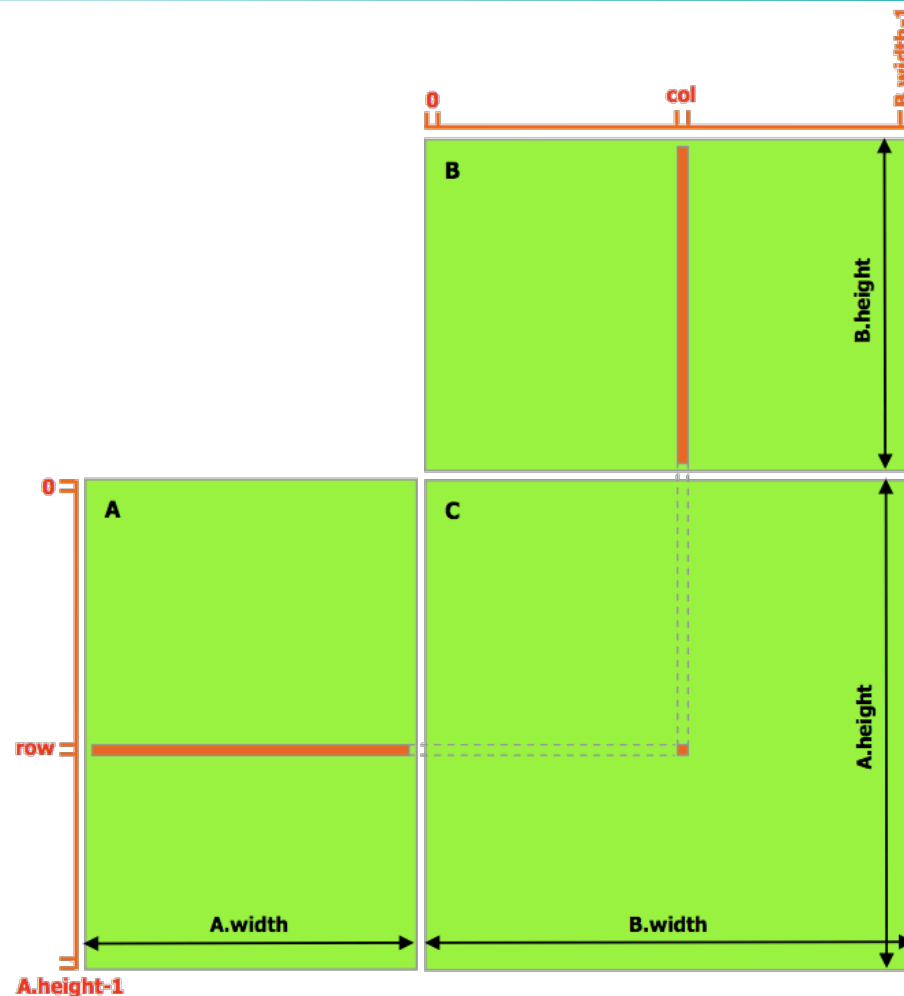
- Platform layer
 - Query and select computing devices
 - Initialize devices
 - Create work contexts and command queues
- Runtime
 - Ressource management
 - Execute kernels
- Compiler
 - Compile and build compute program executable (everything but the host code) at runtime

Outline

- Motivation
- Introducing OpenCL
- Design Goals
- Hardware/ Execution Model
- Software Stack
- Known Example
- Benchmark
- Summary

Known Example

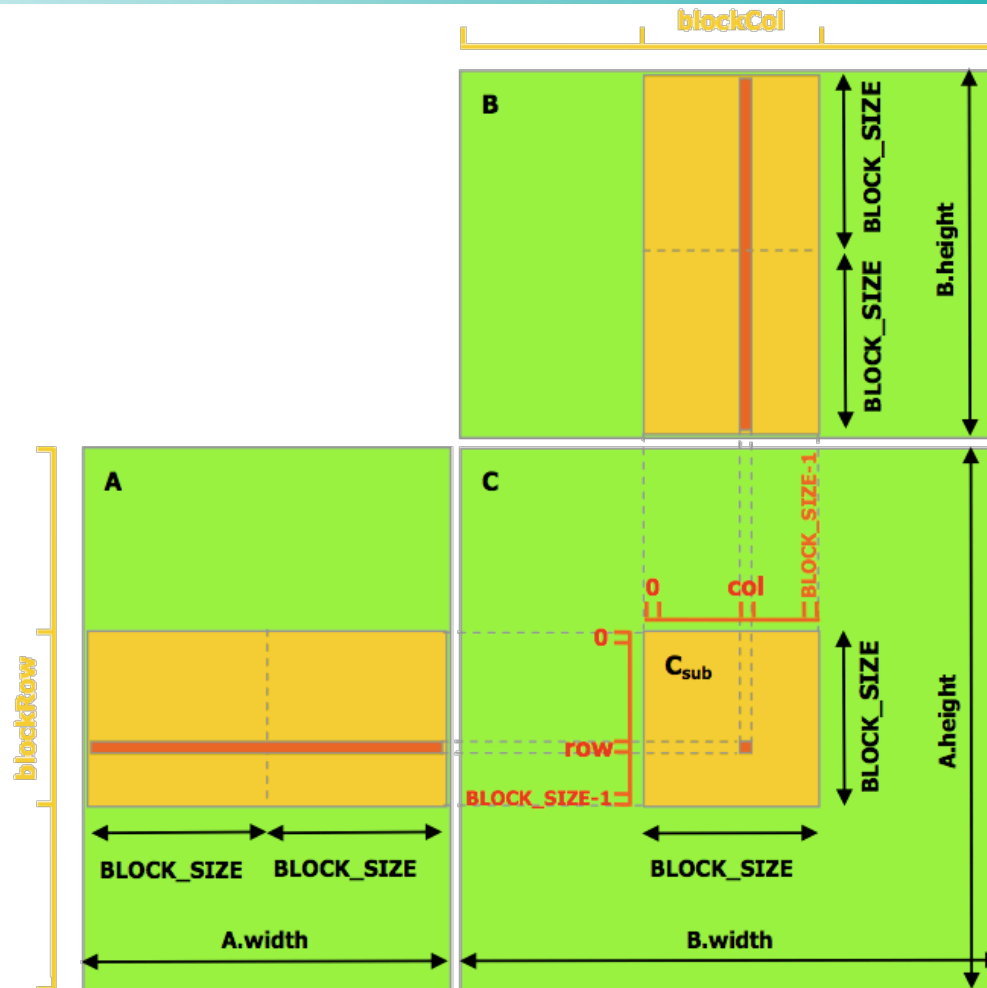
Matrix Multiplication



Source: <http://www.nvidia.com>

Known Example

Matrix Multiplication



Source: <http://www.nvidia.com>

Known Example

Matrix Multiplication

- Code is divided into
 - Host API Code (matrixMul.cpp)
 - OpenCL Kernel Code (matrixMul.cl)
- Kernel code will be familiar to CUDA programmers

Source: <http://www.khronos.org>

Known Example

Compute Kernel Code

```
// Workgroup (thread block) dimensions
#define BLOCK_SIZE 16

// Fast element access
#define AS( i, j) As[ j + i * BLOCK_SIZE]
#define BS( i, j) Bs[ j + i * BLOCK_SIZE]
```


Known Example

Compute Kernel Code

```
__kernel void matrixMul(
    __global float* C, __global float* A, __global float* B,
    __local float* As, __local float* Bs, int widthA, int widthB)
{
    int bx = get_group_id( 0);           // Workgroup (block) index
    int by = get_group_id( 1);           // Workgroup (block) index

    int tx = get_local_id( 0);           // Work-item (thread) index
    int ty = get_local_id( 1);           // Work-item (thread) index

    int aBegin = widthA * BLOCK_SIZE * by; // Sub-matrix A start
    int aEnd   = aBegin + widthA - 1;      // Sub-matrix A end
    int aStep  = BLOCK_SIZE;               // Sub-matrix A stepsize

    int bBegin = BLOCK_SIZE * bx;          // Sub-matrix B start
    int bStep  = BLOCK_SIZE * widthB;      // Sub-matrix B end

    float Csub = 0.0f;                    // Result of block

    // . . .
}
```

matrixMul.cl

Known Example

Compute Kernel Code

```
// . . .
// Loop over A and B to compute the block sub-matrix
for( int a = aBegin, b = bBegin; a <= aEnd; a += aStep, b += bStep)
{
    // Each thread loads one element of the mat. from device to shared memory
    AS( ty, tx) = A[ a + widthA * ty + tx];
    BS( ty, tx) = B[ b + widthB * ty + tx];

    // Synchronize to make sure the matrices are loaded
    barrier( CLK_LOCAL_MEM_FENCE);

    // Multiply the two matrices
    #pragma unroll
    for( int k = 0; k < BLOCK_SIZE; ++k)
        Csub += AS( ty, k) * BS( k, tx);

    // Synchronize to make sure that the preceding
    barrier( CLK_LOCAL_MEM_FENCE);
}
// Write the block sub-matrix to device memory
C[ get_global_id( 1) * get_global_size( 0) + get_global_id( 0)] = Csub;
}
```

matrixMul.cl

Known Example

Host API code

```
#include <CL/opencl.h>

// Matrix dimensions
int WA = 32 * BLOCK_SIZE;           // Input matrix A width
int HA = 128 * BLOCK_SIZE;          // Input matrix A height
int WB = 32 * BLOCK_SIZE;           // Input matrix B width
int HB = WA;                        // Input matrix B height
int WC = WB;                        // Resulting matrix C width
int HC = HA;                        // Resulting matrix C height

int main( int argc, char** argv)
{
    cl_context GPUContext;           // Device context handle
    cl_kernel matMulKernel;          // Kernel handle
    cl_command_queue commandQueue;   // Command queue handle
    cl_program program;              // Program handle
    cl_device_id deviceID = NULL;    // IDs of all OpenCL devices

    // . . .
```

Known Example

Host API code

```
// Allocate memory for matrices on host level
float* hA = (float*) malloc( WA * HA * sizeof(float));
float* hB = (float*) malloc( WB * HB * sizeof(float));
float* hC = (float*) malloc( WC * HC * sizeof(float));

// Fill arrays with random numbers
for(int i = 0; i < WA * HA; i++)
    hA[i] = rand() / (float) RAND_MAX;

for(int i = 0; i < WB * HB; i++)
    hB[i] = rand() / (float) RAND_MAX;

// Get OpenCL devices
clGetDeviceIDs( NULL, CL_DEVICE_TYPE_GPU, 1, &deviceID, NULL);

// Create OpenCL context
GPUContext = clCreateContext( 0, 1, &deviceID, NULL, NULL, NULL);

// Create command queue
commandQueue = clCreateCommandQueue( GPUContext, deviceID,
                                     CL_QUEUE_PROFILING_ENABLE, NULL);

// . . .
```

matrixMul.cpp

Known Example

Host API code

```
// Load OpenCL kernel source code from source file
size_t programLength;
char* source = loadProgSource( SOURCE_PATH, "", &programLength);

// Create OpenCL program
program = clCreateProgramWithSource( GPUContext, 1, (const char**) &source,
                                   &programLength, NULL);

// Build OpenCL program
clBuildProgram( program, 0, NULL, "", NULL, NULL);

// Create OpenCL kernel
matMulKernel = clCreateKernel( program, "matrixMul", NULL);

// . . .
```

matrixMul.cpp

Known Example

Host API code

```
// Create OpenCL buffer pointing to the host memory
cl_mem hABuffer = clCreateBuffer( GPUContext, CL_MEM_READ_ONLY |
                                   CL_MEM_USE_HOST_PTR, WA*HA*sizeof(float), hA, NULL);

// Create buffers (memory on device) for matrices
cl_mem dA = clCreateBuffer( GPUContext, CL_MEM_READ_ONLY,
                             WA*HA*sizeof(float), NULL, NULL);

// Copy from host to device, now
clEnqueueCopyBuffer( commandQueue, hABuffer, dA, 0, 0, WA*HA*sizeof(float),
                    0, NULL, NULL);

// Create buffer on device, it will be initialized from the host at first use
cl_mem dB = clCreateBuffer( GPUContext, CL_MEM_READ_ONLY |
                             CL_MEM_COPY_HOST_PTR, WB*HB*sizeof(float), hB, NULL);

// Create buffer for resulting matrix
cl_mem dC = clCreateBuffer( GPUContext, CL_MEM_WRITE_ONLY,
                             WC*HA*sizeof(float), NULL, NULL);

// . . .
```

matrixMul.cpp

Known Example

Host API code

```
// Set the argument values for the kernel
clSetKernelArg( matMulKernel, 0, sizeof(cl_mem), (void *) &dC);
clSetKernelArg( matMulKernel, 1, sizeof(cl_mem), (void *) &dA);
clSetKernelArg( matMulKernel, 2, sizeof(cl_mem), (void *) &dB);
clSetKernelArg( matMulKernel, 3, BLOCK_SIZE*BLOCK_SIZE*sizeof(float), 0);
clSetKernelArg( matMulKernel, 4, BLOCK_SIZE*BLOCK_SIZE*sizeof(float), 0);
clSetKernelArg( matMulKernel, 5, sizeof(cl_int), (void *) &WA);
clSetKernelArg( matMulKernel, 6, sizeof(cl_int), (void *) &WB);

// Execute Multiplication in parallel
size_t localWorkSize[] = { BLOCK_SIZE, BLOCK_SIZE};
size_t globalWorkSize[] = { WC, HA};

// Launch kernel
// Multiplication – non-blocking execution: launch and push to device
    clEnqueueNDRangeKernel( commandQueue, matMulKernel, 2, 0, globalWorkSize,
        localWorkSize, 0, NULL, NULL);

// Sync to host
clFinish( commandQueue);

// . . .
```

matrixMul.cpp

Known Example

Host API code

```
// Non-blocking copy of result from device to host
clEnqueueReadBuffer( commandQueue, dC, CL_FALSE, 0, WC*HA*sizeof(float),
                    hC, 0, NULL, NULL);

// Release mem objects (device memory)
clReleaseMemObject( hABuffer);
clReleaseMemObject( dA);
clReleaseMemObject( dC);
clReleaseMemObject( dB);

// Clean up OpenCL resources
clReleaseKernel( matMulKernel);
clReleaseCommandQueue( commandQueue);
clReleaseProgram( program);
clReleaseContext( GPUContext);

// Clean up host memory
free( hA); free( hB); free( hC);

// Exit application
return( 0);
}
```

matrixMul.cpp

Outline

- Motivation
- Introducing OpenCL
- Design Goals
- Hardware/ Execution Model
- Software Stack
- Known Example
- Benchmark
- Summary

Benchmark Test System

Intel Core i3-2100

2 cores (2 Threads each) @ 3.1 GHz - 3MB cache - 8GB DDR3 (1333MHz)
max. *52 GFLOPS / s* (LINPACK)

MSI N560GTS-Ti Twin Frozr II/OC

NVIDIA GeForce GTX 560 Ti – 384 cores @ 880Mhz – 1024MB GDDR5 (4200MHz)
max. *1.253 GFLOPS / s* (approximated)

Compiled with Microsoft Visual Studio 2008 Professional x86 under
Microsoft Windows 7 Professional SP1 x64 against

- OpenCL 1.2 (NVIDIA GPU Computing SDK 4.0 x86)
- NVIDIA CUDA 4.0 x86
- NVIDIA CUBLAS 2.0 (NVIDIA GPU Computing SDK 4.0 x86)
- OpenMP 2.0 (Microsoft Visual Studio 2008 x86)

Benchmark

Matrix Multiplication

$$A^{2048 \times 8192} \cdot B^{4096 \times 2048} = C^{4096 \times 8196}$$

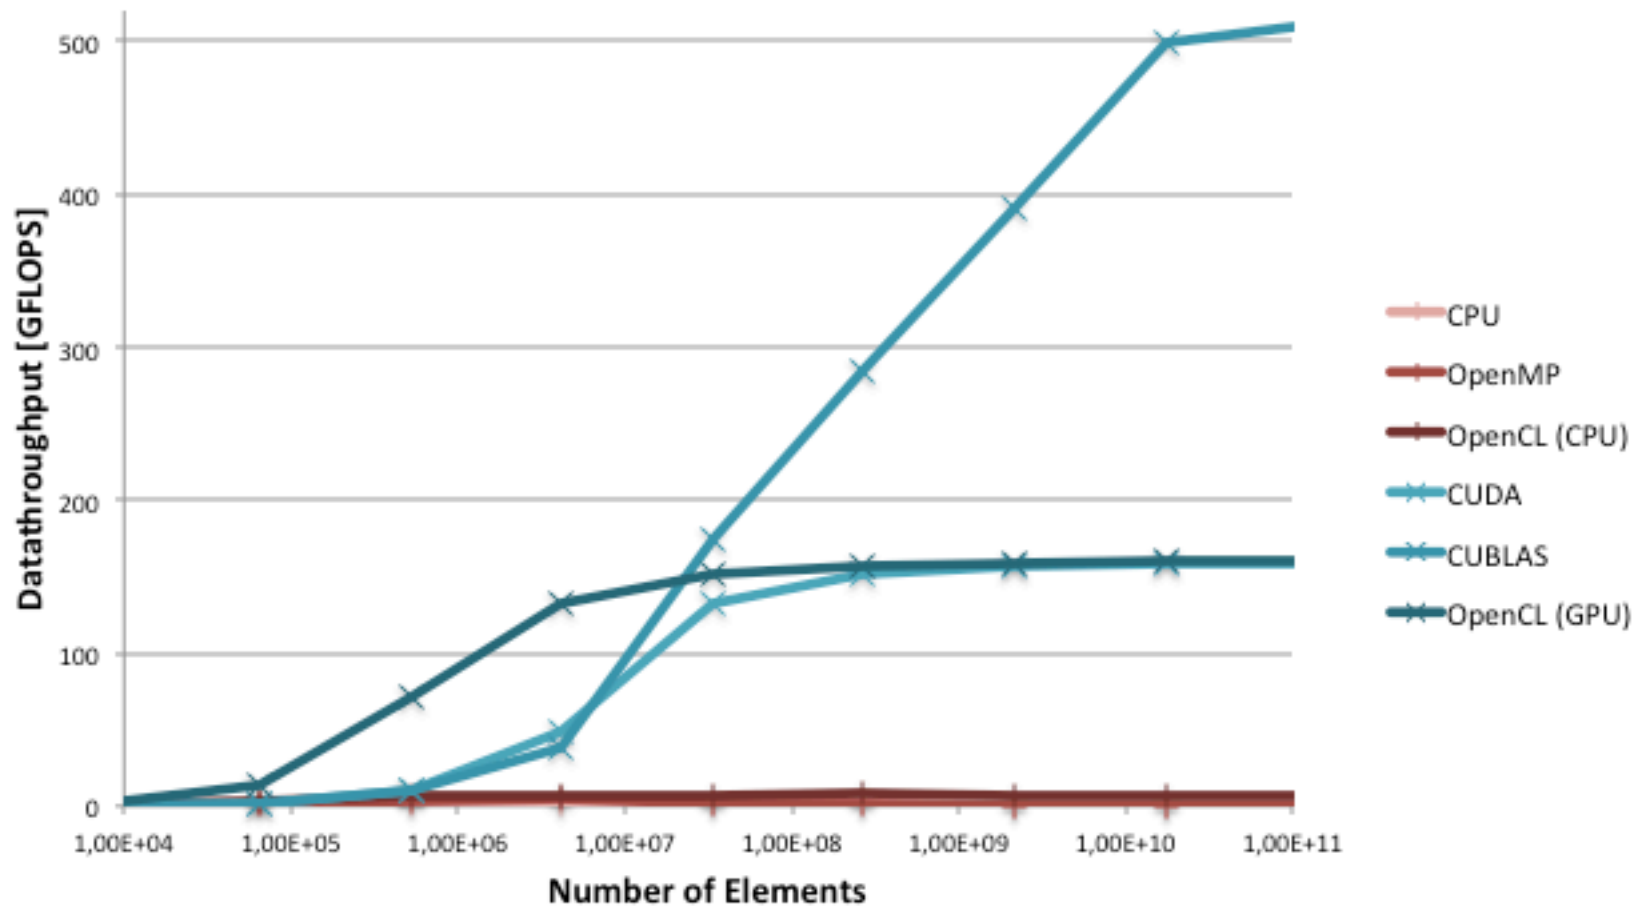
$\Rightarrow 137.4 \cdot 10^9$ elements to compute

| | | Execution Time [s] | Throughput [<i>GFLOPS</i> / s] |
|-----|-------------------|--------------------|---------------------------------|
| CPU | Single threaded | 818.13 | 0.16 |
| | OpenMP | 460.05 | 0.30 |
| | OpenCL 1.2 | 21.76 | 6.32 |
| GPU | NVIDIA CUDA 4.0 | 0.87 | 158.23 |
| | OpenCL 1.2 | 0.86 | 159.72 |
| | NVIDIA CUBLAS 2.0 | 0.24 | 580.90 |

Average after 100 iterations


Benchmark

Matrix Multiplication



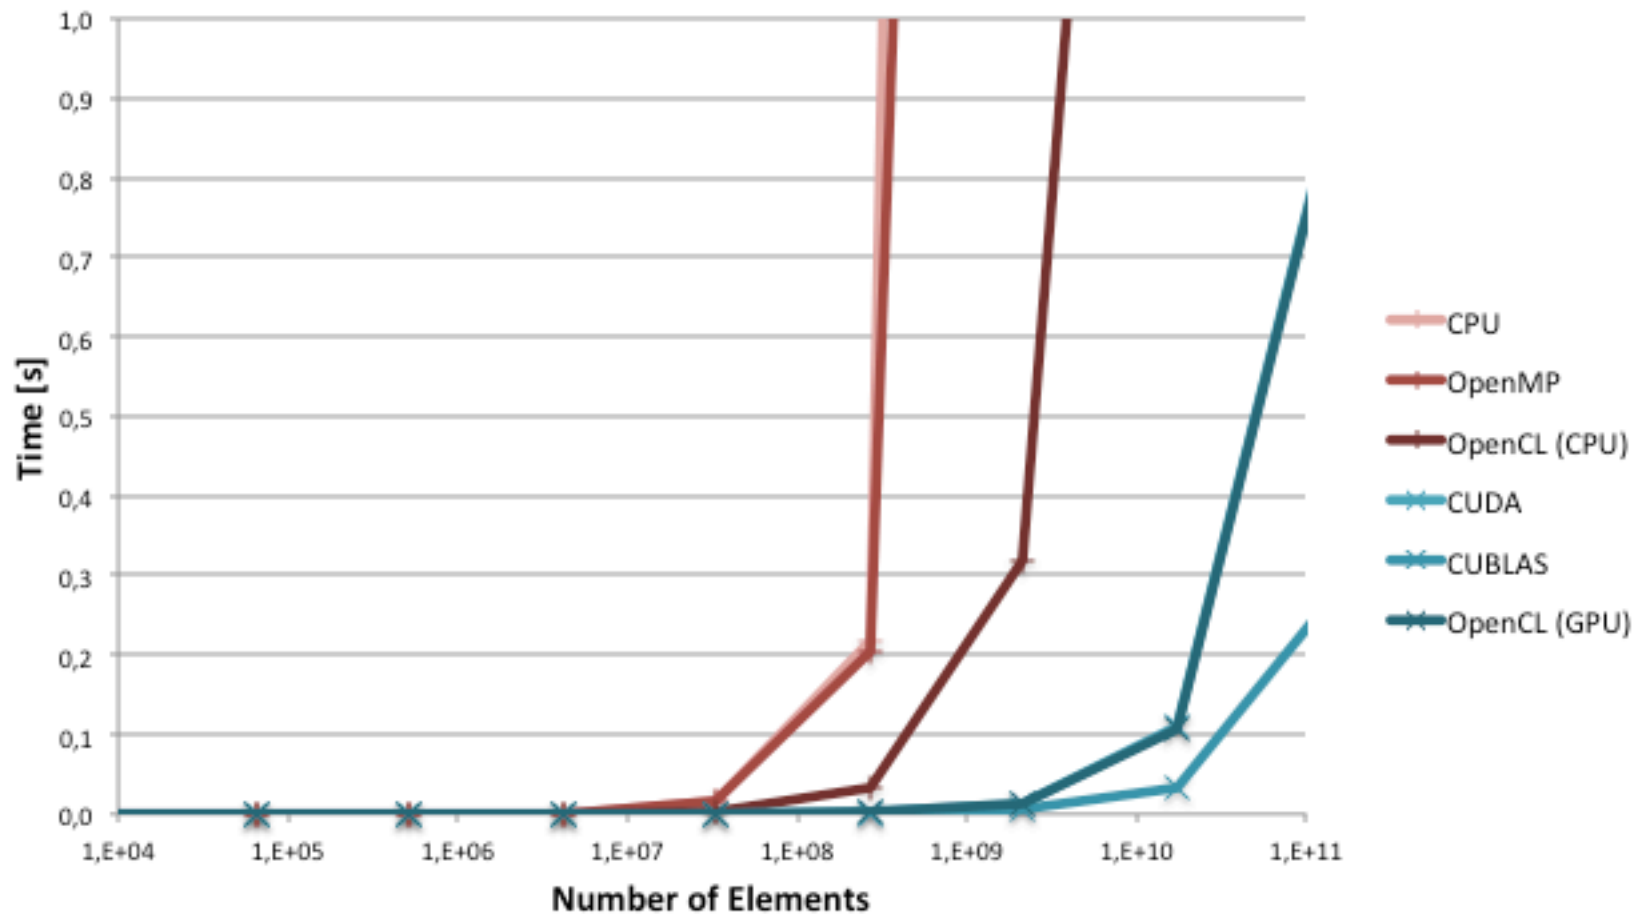
Benchmark

Matrix Multiplication



Benchmark

Matrix Multiplication

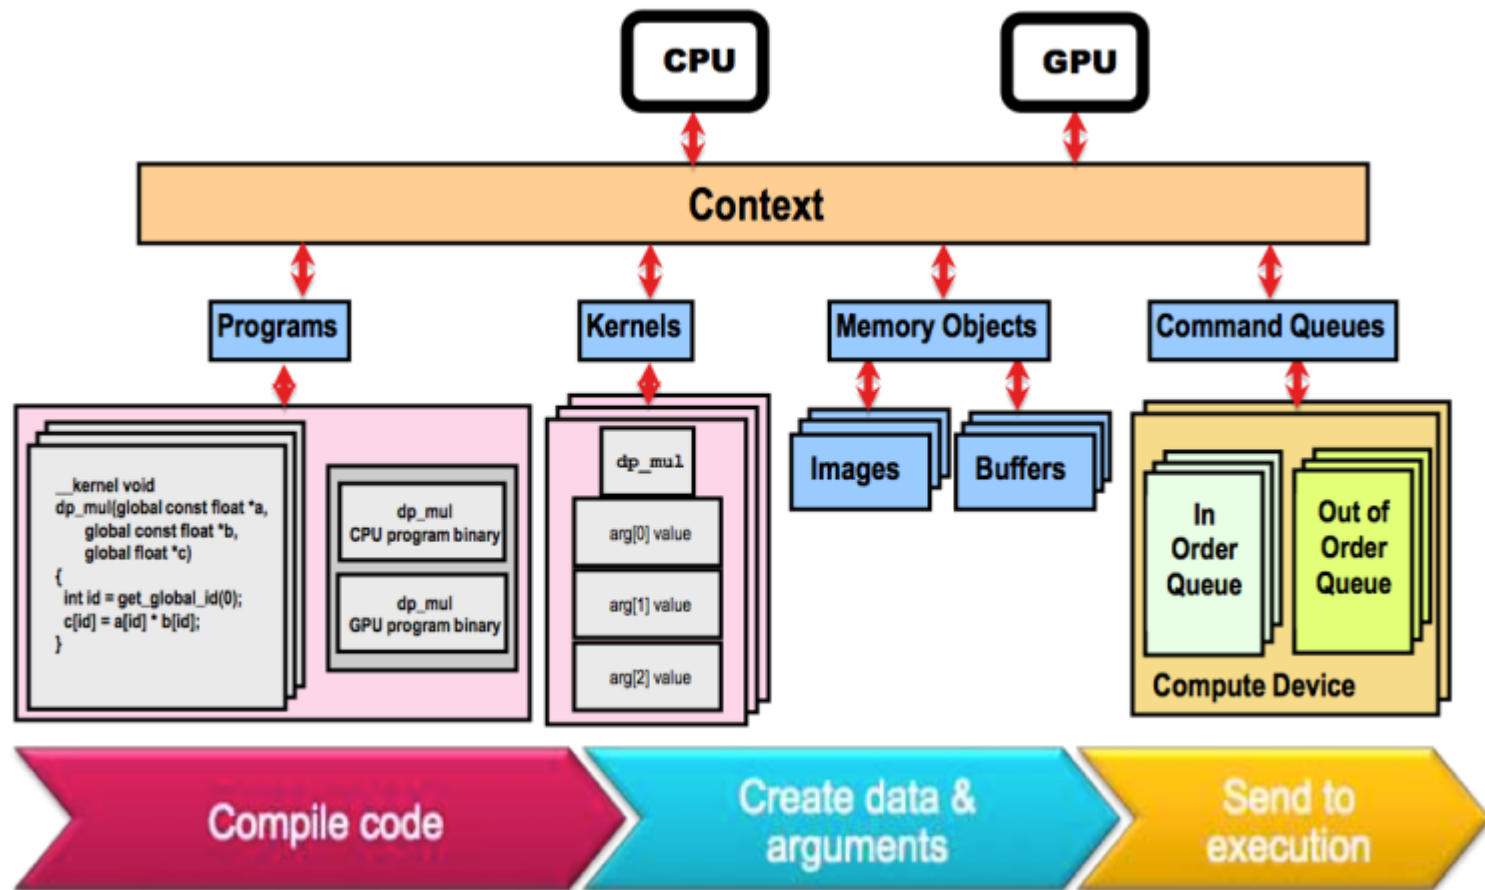


Outline

- Motivation
- Introducing OpenCL
- Design Goals
- Hardware/ Execution Model
- Software Stack
- Known Example
- Benchmark
- Summary

Summary

The OpenCL workflow



Source: <http://www.khronos.org>

Summary

OpenCL vs. CUDA

| | OpenCL 1.2 | NVIDIA CUDA 4.0 |
|---|------------|-----------------|
| Free of charge | ✓ | ✓ |
| Vendor independent/ Multiple SDK-vendors | ✓ | |
| Open standard | ✓ | |
| Use different types of heterogeneous hardware | ✓ | |
| Cross-platform | ✓ | ✓ |
| Interoperability with graphics libraries | ✓ | ✓ |
| Included libraries for common tasks | | ✓ |
| Peer-to-peer-communication between devices | | ✓ |
| Global virtual addressing space | | ✓ |

Summary

Conclusion

OpenCL:

- More complicated to code and redistribute
- Using different types of devices
- Vendor independend
- Outstanding performance on many CPUs

NVIDIA CUDA:

- More enhanced
- Already widely spread
- Deliveres well-engineered libraries for many tasks

Thank you for your attention.

Source code at <http://janbeneke.de/ParComp>

References

- Khronos Group khronos.org/opencv/
- NVIDIA developer.nvidia.com/opencv/
- Intel software.intel.com/en-us/articles/vcsource-tools-opencv-sdk/
- AMD/ ATI developer.amd.com/zones/OpenCLZone/